

Functions

STAT 133

Gaston Sanchez

Department of Statistics, UC–Berkeley

`gastonsanchez.com`

`github.com/gastonstat/stat133`

Course web: gastonsanchez.com/teaching/stat133

Functions

R comes with many functions and packages that let us perform a wide variety of tasks. But there are occasions in which we need to create our own functions.

Anatomy of a function

`function()` allows us to create a function. It has the following structure:

```
function_name <- function(arg1, arg2, ...)  
{  
  expressions  
}
```

Anatomy of a function

- ▶ Generally we will give a name to a function
- ▶ A function takes one or more arguments (or none)
- ▶ The expressions forming the operations comprise the body of the function
- ▶ Braces surround the body of the function
- ▶ Functions return a single *value*

Function example

A function that squares its argument:

```
square <- function(x) {  
  x * x  
}
```

It works like any other function in R:

```
square(10)
```

```
## [1] 100
```

In this case, `square()` is also vectorized

```
square(1:5)
```

```
## [1] 1 4 9 16 25
```

Function example

Functions with a body consisting of a simple expression can be written with no braces (in one single line!):

```
square <- function(x) x * x
```

```
square(10)
```

```
## [1] 100
```

Function example

Once defined, functions can be used in other function definitions:

```
sum_square <- function(x) sum(square(x))
```

```
sum_square(1:5)
```

```
## [1] 55
```

A simple example

A function which, given the value r computes the value πr^2

```
area <- function(r) pi * r^2
```

- ▶ The formal argument of the function is r
- ▶ The body of the function consists of the simple expression $\pi * r^2$
- ▶ The function has been assigned the name "area"

```
area(5)
```

```
## [1] 78.53982
```


Evaluation of Functions

Function evaluation involves:

- ▶ A set of variables associated to the arguments is temporarily created
- ▶ The variable definitions are used to evaluate the body function
- ▶ Temporary variables are removed at the end
- ▶ The computed values are returned

Evaluation Example

Evaluating the function call `area(5)` takes place as follows:

- ▶ Temporarily create a variable `r` with value 5
- ▶ Use that value 5 to compute `pi * 5^2`
- ▶ Remove the temporary variable definition
- ▶ Return the value 78.53982

Function example

```
hello <- function(x) {  
  paste("Hello", x)  
}
```

```
hello('Gaston')
```

```
## [1] "Hello Gaston"
```

Another function example

```
add <- function(x, y) {  
  x + y  
}
```

```
add(2, 3)
```

```
## [1] 5
```

Function with no arguments

Functions can have no arguments

```
hi <- function() {  
  print("Hi there!")  
}
```

```
hi()
```

```
## [1] "Hi there!"
```

Missing arguments

If you specify an argument with no default value, you must give it a value everytime you call the function, otherwise you'll get an error:

```
sqr <- function(x) {  
  x^2  
}
```

```
sqr()
```

```
## Error in sqr(): argument "x" is missing, with no  
default
```

Default arguments

You can give default values to function arguments:

```
hey <- function(x = "") {  
  cat("Hey", x, "\nHow is it going?" )  
}
```

```
hey()
```

```
## Hey  
## How is it going?
```

```
hey("Gaston")
```

```
## Hey Gaston  
## How is it going?
```

The return() command

Sometimes the `return()` command is included to explicitly indicate the output of a function:

```
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}
```

```
add(2, 3)
```

```
## [1] 5
```


The return() command

If no return() is present, then R returns the last evaluated expression:

```
# output with return()  
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}  
  
add(2, 3)  
  
## [1] 5
```

```
# output without return()  
add <- function(x, y) {  
  x + y  
}  
  
add(2, 3)  
  
## [1] 5
```

The return() command

If no return() is present, then R returns the last evaluated expression, although it might not always print the output:

```
# nothing is printed  
add <- function(x, y) {  
  z <- x + y  
}  
  
add(2, 3)
```

```
# output printed  
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}  
  
add(2, 3)  
  
## [1] 5
```

The return() command

The last evaluated expression has the same value in both cases:

```
# nothing is printed
add <- function(x, y) {
  z <- x + y
}

a1 <- add(2, 3)
a1

## [1] 5
```

```
# output printed
add <- function(x, y) {
  z <- x + y
  return(z)
}

a2 <- add(2, 3)
a2

## [1] 5
```

The return() command

If no return() is present, then R returns the last evaluated expression:

```
add1 <- function(x, y) {  
  x + y  
}
```

```
add2 <- function(x, y) {  
  z <- x + y  
  z  
}
```

```
add3 <- function(x, y) {  
  z <- x + y  
}
```

```
add4 <- function(x, y) {  
  z <- x + y  
  return(z)  
}
```

The return() command

return() can be useful when the output may be obtained in the middle of the function's body

```
f <- function(x, y, add = TRUE) {  
  if (add) {  
    return(x + y)  
  } else {  
    return(x - y)  
  }  
}
```

```
f(2, 3, add = TRUE)
```

```
## [1] 5
```

```
f(2, 3, add = FALSE)
```

```
## [1] -1
```

Exercise

Write a function that checks if a number is positive (output TRUE) or negative (output FALSE)

Exercise

Write a function that checks if a number is positive (output TRUE) or negative (output FALSE)

```
is_positive <- function(x) {  
  if (x > 0) TRUE else FALSE  
}
```

```
is_positive(2)
```

```
## [1] TRUE
```

```
is_positive(-1)
```

```
## [1] FALSE
```

Exercise

What happens in these cases?

```
is_positive <- function(x) {  
  if (x > 0) TRUE else FALSE  
}
```

```
is_positive(0)
```

```
is_positive(NA)
```

```
is_positive(TRUE)
```

```
is_positive("positive")
```

```
is_positive(1:5)
```


Writing Functions

Writing Functions

Writing Functions

- ▶ Choose meaningful names of functions
- ▶ Preferably a verb
- ▶ Choose meaningful names of arguments
- ▶ Think about the users (who will use the function)
- ▶ Think about extreme cases
- ▶ If a function is too long, maybe you need to split it

Names of functions

Avoid this:

```
f <- function(x, y) {  
  x + y  
}
```

This is better

```
add <- function(x, y) {  
  x + y  
}
```

Describing functions

Also add a short description of what the arguments should be like. In this case, the description is outside the function

```
# function for adding two numbers  
# x: number  
# y: number  
add <- function(x, y) {  
  x + y  
}
```

Describing functions

In this case, the description is inside the function

```
add <- function(x, y) {  
  # function for adding two numbers  
  # x: number  
  # y: number  
  x + y  
}
```

Exercise

R has a function `summary()` that when applied on a numeric vector provides something like this:

```
summary(1:10)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.00	3.25	5.50	5.50	7.75	10.00

Create a `describe()` function that takes a numeric vector and returns: minimum, maximum, mean, and standard deviation

Exercise

First attempt

```
describe <- function(x) {  
  x_min <- min(x)  
  x_max <- max(x)  
  x_mean <- mean(x)  
  x_sd <- sd(x)  
  return(c(x_min, x_max, x_mean, x_sd))  
}  
  
describe(1:10)  
  
## [1] 1.00000 10.00000 5.50000 3.02765
```

Exercise

Second attempt (adding names)

```
describe <- function(x) {  
  x_min <- min(x)  
  x_max <- max(x)  
  x_mean <- mean(x)  
  x_sd <- sd(x)  
  values <- c(x_min, x_max, x_mean, x_sd)  
  names(values) <- c("min", "max", "mean", "sd")  
  return(values)  
}
```

```
describe(1:10)
```

```
##      min      max      mean      sd  
## 1.00000 10.00000  5.50000  3.02765
```


Exercise

Third attempt (using a list as output)

```
describe <- function(x) {  
  list(  
    min = min(x),  
    max = max(x),  
    mean = mean(x),  
    sd = sd(x)  
  )  
}
```

```
describe(1:10)  
  
## $min  
## [1] 1  
##  
## $max  
## [1] 10  
##  
## $mean  
## [1] 5.5  
##  
## $sd  
## [1] 3.02765
```

Handling Function Arguments

Exercise

Assertions

```
describe <- function(x) {  
  if (!is.numeric(x)) {  
    stop("x is not numeric")  
  }  
  # output  
  c(mean = mean(x), sd = sd(x))  
}
```

Exercise

Passing arguments

```
describe <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    stop("x is not numeric")  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

```
describe(c(1:5, NA), na.rm = TRUE)
```

```
##      mean      sd  
## 3.000000 1.581139
```

Exercise

Probability Density of the Normal Distribution:

$$f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Write a function that takes a value x (with parameters μ and σ) which computes the probability density distribution of the normal distribution

Exercise

Normal Distribution:

```
normal_dist <- function(x, mu = 0, sigma = 1) {  
  constant <- 1 / (sigma * sqrt(2*pi))  
  constant * exp(-((x - mu)^2) / (2 * sigma^2))  
}  
  
normal_dist(2)  
  
## [1] 0.05399097
```

Argument Matching

- ▶ Arguments can be named or not
- ▶ Named arguments are used in preference to position which
- ▶ It is important to be clear about which argument corresponds to which formal parameter of the function

Argument Matching

```
normal_dist <- function(x, mu = 0, sigma = 1) {  
  constant <- 1 / (sigma * sqrt(2*pi))  
  constant * exp(-((x - mu)^2) / (2 * sigma^2))  
}
```

```
normal_dist(2)  
normal_dist(2, sigma = 3, mu = 1)  
normal_dist(mu = 1, sigma = 3, 2)  
normal_dist(mu = 1, 2, sigma = 3)
```

Argument Matching

R is “smart” enough in doing pattern matching with arguments' names (not recommended though)

```
normal_dist(2)
```

```
## [1] 0.05399097
```

```
normal_dist(2, m = 0, s = 1)
```

```
## [1] 0.05399097
```

```
normal_dist(2, sig = 1, m = 0)
```

```
## [1] 0.05399097
```