

Functions

STAT 133

Gaston Sanchez

Department of Statistics, UC–Berkeley

`gastonsanchez.com`

`github.com/gastonstat/stat133`

Course web: `gastonsanchez.com/teaching/stat133`

Functions

R comes with many functions and packages that let us perform a wide variety of tasks. Sometimes, however, there's no function to do what we want to achieve. In these cases we need to create our own functions.

Anatomy of a Function

Anatomy of a function

`function()` allows us to create a function. It has the following structure:

```
function_name <- function(arg1, arg2, etc)
{
  expression_1
  expression_2
  ...
  expression_n
}
```

Anatomy of a function

- ▶ Generally we will give a name to a function
- ▶ A function takes one or more inputs (or none), known as *arguments*
- ▶ The expressions forming the operations comprise the body of the function
- ▶ Simple expression don't require braces
- ▶ Compound expressions are surround by braces
- ▶ Functions return a single *value*

Function example

A function that squares its argument:

```
square <- function(x) {  
  x * x  
}
```

- ▶ the function's name is "square"
- ▶ it has one argument `x`
- ▶ the function's body consists of one simple expression
- ▶ it returns the value `x * x`

Function example

It works like any other function in R:

```
square(10)
```

```
## [1] 100
```

In this case, `square()` is also vectorized

```
square(1:5)
```

```
## [1] 1 4 9 16 25
```

Why is `square()` vectorized?

Function example

Functions with a body consisting of a simple expression can be written with no braces (in one single line!):

```
square <- function(x) x * x
```

```
square(10)
```

```
## [1] 100
```


If the body of a function is a compound expression we use braces:

```
sum_sqr <- function(x, y) {  
  xy_sum <- x + y  
  xy_ssqr <- (xy_sum)^2  
  list(sum = xy_sum,  
        sumsqr = xy_ssqr)  
}
```

```
sum_sqr(3, 5)
```

```
## $sum  
## [1] 8  
##  
## $sumsqr  
## [1] 64
```

Function example

Once defined, functions can be used in other function definitions:

```
sum_square <- function(x) sum(square(x))
```

```
sum_square(1:5)
```

```
## [1] 55
```

Area of a Circle

A function which, given the value r computes the value πr^2

```
area_circle <- function(r) pi * r^2
```

- ▶ The formal argument of the function is r
- ▶ The body of the function consists of the simple expression $\text{pi} * r^2$
- ▶ The function has been assigned the name "area_circle"

```
area_circle(5)
```

```
## [1] 78.53982
```

Evaluation of Functions

Function evaluation involves:

- ▶ A set of variables associated to the arguments is temporarily created
- ▶ The variable definitions are used to evaluate the body function
- ▶ Temporary variables are removed at the end
- ▶ The computed values are returned

Evaluation Example

Evaluating the function call `area_circle(5)` takes place as follows:

- ▶ Temporarily create a variable `r` with value 5
- ▶ Use that value 5 to compute `pi * 5^2`
- ▶ Remove the temporary variable definition
- ▶ Return the value `78.53982`

Nested Functions

We can also define a function inside another function:

```
getmax <- function(a) {  
  maxpos <- function(u) which.max(u)  
  list(position = maxpos(a),  
        value = max(a))  
}
```

```
getmax(c(2, -4, 6, 10, pi))
```

```
## $position  
## [1] 4  
##  
## $value  
## [1] 10
```

Function names

Different ways to name functions

- ▶ `squareroot()`
- ▶ `SquareRoot()`
- ▶ `squareRoot()`
- ▶ `square.root()`
- ▶ `square_root()`

Function names

Invalid names

- ▶ `5squareroot()`: cannot begin with a number
- ▶ `_sqrt()`: cannot begin with an underscore
- ▶ `square-root()`: cannot use hyphenated names

In addition, avoid using an already existing name, e.g. `sqrt()`

Function Output

Function output

- ▶ The body of a function is an expression
- ▶ Remember that every expression has a value
- ▶ Hence every function has a value

Function output

The value of a function can be established in two ways:

- ▶ As the last evaluated simple expression (in the body)
- ▶ An explicitly **returned** value via `return()`

The return() command

Sometimes the `return()` command is included to explicitly indicate the output of a function:

```
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}
```

```
add(2, 3)
```

```
## [1] 5
```

The return() command

If no return() is present, then R returns the last evaluated expression:

```
# output with return()  
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}  
  
add(2, 3)  
  
## [1] 5
```

```
# output without return()  
add <- function(x, y) {  
  x + y  
}  
  
add(2, 3)  
  
## [1] 5
```

The return() command

Depending on what's returned or what's the last evaluated expression, just calling a function might not print anything:

```
# nothing is printed  
add <- function(x, y) {  
  z <- x + y  
}  
  
add(2, 3)
```

```
# output printed  
add <- function(x, y) {  
  z <- x + y  
  return(z)  
}  
  
add(2, 3)  
  
## [1] 5
```

The return() command

Here we call the function and assign it to an object. The last evaluated expression has the same value in both cases:

```
# nothing is printed
add <- function(x, y) {
  z <- x + y
}

a1 <- add(2, 3)
a1

## [1] 5
```

```
# output printed
add <- function(x, y) {
  z <- x + y
  return(z)
}

a2 <- add(2, 3)
a2

## [1] 5
```

The return() command

If no return() is present, then R returns the last evaluated expression:

```
add1 <- function(x, y) {  
  x + y  
}
```

```
add2 <- function(x, y) {  
  z <- x + y  
  z  
}
```

```
add3 <- function(x, y) {  
  z <- x + y  
}
```

```
add4 <- function(x, y) {  
  z <- x + y  
  return(z)  
}
```


The return() command

return() can be useful when the output may be obtained in the middle of the function's body

```
f <- function(x, y, add = TRUE) {  
  if (add) {  
    return(x + y)  
  } else {  
    return(x - y)  
  }  
}
```

```
f(2, 3, add = TRUE)  
## [1] 5  
  
f(2, 3, add = FALSE)  
## [1] -1
```

Function Arguments

Function arguments

Functions can have any number of arguments (even zero arguments)

```
# function with 2 arguments  
add <- function(x, y) x + y  
  
# function with no arguments  
hi <- function() print("Hi there!")  
  
hi()  
  
## [1] "Hi there!"
```

Arguments

Arguments can have default values

```
hey <- function(x = "") {  
  cat("Hey", x, "\nHow is it going?" )  
}
```

```
hey()
```

```
## Hey  
## How is it going?
```

```
hey("Gaston")
```

```
## Hey Gaston  
## How is it going?
```

Arguments with no default values

If you specify an argument with no default value, you must give it a value everytime you call the function, otherwise you'll get an error:

```
sqr <- function(x) {  
  x^2  
}
```

```
sqr()
```

```
## Error in sqr(): argument "x" is missing, with no  
default
```

Arguments with no default values

Sometimes we don't want to give default values, but we also don't want to cause an error. We can use `missing()` to see if an argument is missing:

```
abc <- function(a, b, c = 3) {  
  if (missing(b)) {  
    result <- a * 2 + c  
  } else {  
    result <- a * b + c  
  }  
  result  
}
```

```
abc(1)
```

```
## [1] 5
```

```
abc(1, 4)
```

```
## [1] 7
```

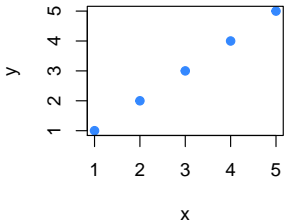
Arguments with no default values

You can also set an argument value to NULL if you don't want to specify a default value:

```
abcd <- function(a, b = 2, c = 3, d = NULL) {  
  if (is.null(d)) {  
    result <- a * b + c  
  } else {  
    result <- a * b + c * d  
  }  
  result  
}
```

More arguments

```
# arguments with and without default values  
myplot <- function(x, y, col = "#3488ff", pch = 19) {  
  plot(x, y, col = col, pch = pch)  
}  
  
myplot(1:5, 1:5)
```



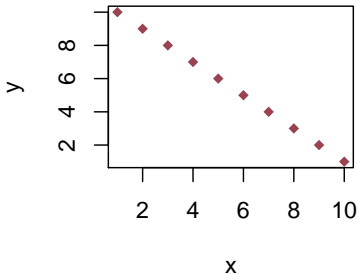
More arguments

```
# arguments with and without default values
myplot <- function(x, y, col = "#3488ff", pch = 19) {
  plot(x, y, col = col, pch = pch)
}
```

- ▶ x and y have no default values
- ▶ col and pch have default values (but they can be changed)

More arguments

```
# changing default arguments  
myplot(1:10, 10:1, col = "#994352", pch = 18)
```



Argument Matching

- ▶ Arguments with default values are known as **named** arguments
- ▶ Arguments with no default values are referred to as **positional** arguments
- ▶ Arguments can be matched positionally or by name

Argument Matching

```
values <- seq(-2, 1, length.out = 20)
```

```
# equivalent calls
```

```
mean(values)
```

```
mean(x = values)
```

```
mean(x = values, na.rm = FALSE)
```

```
mean(na.rm = FALSE, x = values)
```

```
mean(na.rm = FALSE, values)
```

Partial Matching

Named arguments can also be partially matched:

```
# equivalent calls  
seq(from = 1, to = 2, length.out = 5)  
seq(from = 1, to = 2, length = 5)  
seq(from = 1, to = 2, len = 5)
```

`length.out` is partially matched with `length` and `len`

Matching Order

Order of argument matching operations:

- ▶ Check for exact match for a named argument
- ▶ Check for a partial match
- ▶ Check for a positional match

Exercise

Write a function that checks if a number is positive (output TRUE) or negative (output FALSE)

Exercise

Write a function that checks if a number is positive (output TRUE) or negative (output FALSE)

```
is_positive <- function(x) {  
  if (x > 0) TRUE else FALSE  
}
```

```
is_positive(2)
```

```
## [1] TRUE
```

```
is_positive(-1)
```

```
## [1] FALSE
```


Exercise

What happens in these cases?

```
is_positive <- function(x) {  
  if (x > 0) TRUE else FALSE  
}
```

```
is_positive(0)
```

```
is_positive(NA)
```

```
is_positive(TRUE)
```

```
is_positive("positive")
```

```
is_positive(1:5)
```

Using arguments for other functions

There are various functions that include the argument `na.rm` to indicate if missing values should be removed. One of them is `mean()`:

```
# default na.rm = FALSE
mean(c(1, 2, 3, NA, 5))

## [1] NA

# default na.rm = TRUE
mean(c(1, 2, 3, NA, 5), na.rm = TRUE)

## [1] 2.75
```

Using arguments for other functions

If we create a function that uses other functions containing `na.rm`, it is wise to include that argument:

```
meansd <- function(x, na.rm = FALSE) {  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}  
  
meansd(c(1, 2, 3, NA, 5), na.rm = TRUE)  
  
##      mean      sd  
## 2.750000 1.707825
```

Dots argument

If you check functions like `c()`, `paste()`, `plot()`, you'll notice the use of a special argument: `...`

- ▶ it matches zero, one or more actual arguments
- ▶ it allows us to pass arguments to other functions inside the function
- ▶ `...` allows us to “cascade” arguments to other functions without including them in the definition

Dots argument

Using ...

```
fplot <- function(y, ...) {  
  x <- 1:length(y)  
  plot(x, y, type = 'n', ylim = c(0, 1), ...)  
  points(x, y, col = "#93a8f2", pch = 19)  
  lines(x, y, col = "#dd93f2", lwd = 3)  
}  
  
fplot(runif(10), bty = 'n')  
fplot(runif(10), bty = 'n', main = "some title")  
fplot(runif(10), bty = 'n', xlab = '')
```

Writing Functions

Writing Functions

Writing Functions

- ▶ Choose meaningful names of functions
- ▶ Preferably a verb
- ▶ Choose meaningful names of arguments
- ▶ Think about the users (who will use the function)
- ▶ Think about extreme cases
- ▶ If a function is too long, maybe you need to split it

Names of functions

Avoid this:

```
f <- function(x, y) {  
  x + y  
}
```

This is better

```
add <- function(x, y) {  
  x + y  
}
```


Describing functions

Also add a short description of what the arguments should be like. In this case, the description is outside the function

```
# function for adding two numbers  
# x: number  
# y: number  
add <- function(x, y) {  
  x + y  
}
```

Describing functions

In this case, the description is inside the function

```
add <- function(x, y) {  
  # function for adding two numbers  
  # x: number  
  # y: number  
  x + y  
}
```

Binary Operators

Binary Operators

- ▶ One type of functions very common in R are **binary operators**, eg:
 - $2 + 5$ (sum)
 - $3 / 2$ (exponent)
 - $a \%in\% b$ (value matching)
 - $X \%*\% Y$ (matrix multiplication)
- ▶ These functions take two inputs—hence the term *binary*
- ▶ It is possible to define your own binary operators

Binary Operators

```
# binary operator
"%p%" <- function(x, y) {
  paste(x, y, sep = " ")
}

'good' %p% 'morning'

## [1] "good morning"
```

How to create a binary operator?

- ▶ A binary operator is defined as one or more characters surrounded by percent symbols %
- ▶ When defining the function, the entire name must be quoted
- ▶ Include two arguments
- ▶ As usual, avoid using names of existing operators:
 - "%%", "%*", "%/%", "%o%", "%in%

Another example

Here's another example:

```
# binary operator
"%u%" <- function(x, y) {
  union(x, y)
}
```

```
1:5 %u% c(1, 3, 5, 7, 9)
```

```
## [1] 1 2 3 4 5 7 9
```

Lazy Evaluation

Lazy Evaluation

Arguments to functions are evaluated lazily, that is, they are evaluated only as needed:

```
g <- function(a, b) {  
  a * a * a  
}
```

```
g(2)
```

```
## [1] 8
```

`g()` never uses the argument `b`, so calling `g(2)` does not produce an error

Lazy Evaluation

Another example

```
g <- function(a, b) {  
  print(a)  
  print(b)  
}
```

```
g(2)
```

```
## [1] 2
```

```
## Error in print(b): argument "b" is missing, with no  
default
```

Notice that 2 got printed before the error was triggered. This is because b did not have to be evaluated until after print(a)

Messages

There are two main functions for generating warnings and errors:

- ▶ `stop()`
- ▶ `warning()`

There's also the `stopifnot()` function

Stop Execution

Use `stop()` to stop the execution of a function (this will raise an error)

```
meansd <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    stop("x is not numeric")  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

Stop Execution

Use `stop()` to stop the execution of a function (this will raise an error)

```
# ok  
meansd(c(4, 5, 3, 1, 2))
```

```
##      mean      sd  
## 3.000000 1.581139
```

```
# this causes an error  
meansd(c('a', 'b', 'c'))
```

```
## Error in meansd(c("a", "b", "c")): x is not numeric
```

Warning Messages

Use `warning()` to show a warning message

```
meansd <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    warning("non-numeric input coerced to numeric")  
    x <- as.numeric(x)  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

A warning is useful when we don't want to stop the execution, but we still want to show potential problems

Warning Messages

Use `warning()` to show a warning message

```
# ok
meansd(c(4, 5, 3, 1, 2))

##          mean          sd
## 3.000000 1.581139

# this causes a warning
meansd(c(TRUE, FALSE, TRUE, FALSE))

## Warning in meansd(c(TRUE, FALSE, TRUE, FALSE)):
## non-numeric input coerced to numeric

##          mean          sd
## 0.5000000 0.5773503
```


Stop Execution

`stopifnot()` ensures the truth of expressions:

```
meansd <- function(x, na.rm = FALSE) {  
  stopifnot(is.numeric(x))  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
      sd = sd(x, na.rm = na.rm))  
}
```

```
meansd('hello')
```

```
## Error:  is.numeric(x) is not TRUE
```

Environments and Functions

Consider this example

```
w <- 10

f <- function(y) {
  d <- 5
  h <- function() {
    d * (w + y)
  }
  return(h())
}

f(2)

## [1] 60
```

How / Why does `f()` work?

Consider this other example

```
w <- 10

f <- function(y) {
  d <- 5
  return(h())
}

f(2)

## Error in f(2):  could not find function "h"
```

Why `f()` does not work?

Environments

- ▶ All the variables that we create need to be stored somewhere
- ▶ The place where they are stored is called an **environment**
- ▶ R works with environments, all of which are in (virtual) memory
- ▶ Usually, we don't need to explicitly deal with environments
- ▶ Environments are nested

Global Environment

- ▶ The user workspace is the **global environment**
- ▶ The global environment is the **top level** environment
- ▶ It is formally referred to as `R_GlobalEnv`
- ▶ Variables defined in the global environment can be seen from anywhere
- ▶ The contents of the global environment are listed with `ls()`

```
# top level environment  
environment()
```

```
## <environment: R_GlobalEnv>
```

Searching objects

- ▶ When R tries to bind a value to a symbol, it searches through a series of environments to find the appropriate value
- ▶ To retrieve the value of an object the order is:
- ▶ Search the current environment
- ▶ Search the global environment for a symbol name matching the one requested
- ▶ Search the namespaces of each of the packages on the search list: `search()`

Environments and Functions

- ▶ A function consists not only of its arguments and body but also of its *environment*
- ▶ An environment is made up of the collection of objects present at the time the function comes into existence
- ▶ When a function is created by evaluating the corresponding expression, the current environment is recorded as a property of the function

Let's go back to our first example

```
w <- 10

f <- function(y) {
  d <- 5
  h <- function() {
    d * (w + y)
  }
  return(h())
}

f(2)

## [1] 60
```

How does `f()` work?

Let's see the environments

```
w <- 10    # variable (in global environment)

# a function (in global environment)
f <- function(y) {
  d <- 5    # local variable
  h <- function() {    # subfunction
    d * (w + y)    # w is a free variable
  }
  return(h())
}

environment(f)

## <environment: R_GlobalEnv>
```

Function Environment

- ▶ `w` is a global variable (in global environment)
- ▶ `f()` is a function in the global environment
- ▶ `d` is a local variable—local to `f()`
- ▶ `h()` is a subfunction—local to `f()`
- ▶ `w` is not an argument but a free variable

Let's see the environments

```
f <- function(y) {  
  d <- 5  
  h <- function() {  
    d * (w + y)  
  }  
  print(environment(h))    # h()'s environment  
  return(h())  
}  
  
environment(f)  
  
## <environment: R_GlobalEnv>  
  
f(2)  
  
## <environment: 0x7fe39b35aea0>  
## [1] 60
```

Variable's Scope

- ▶ A variable's **scope** is the set of places from which you can see the variable
- ▶ R will try to find a variable in the current environment
- ▶ If it doesn't find them it will look in the parent environment
- ▶ And then that environment's parent
- ▶ And so on until it reaches the global environment

Variable's Scope

- ▶ A variable's **scope** is the set of places from which you can see the variable
- ▶ R will try to find a variable in the current environment
- ▶ If it doesn't find them it will look in the parent environment
- ▶ And then that environment's parent
- ▶ And so on until it reaches the global environment

Variable Scope

- ▶ When we define a variable inside a function, the rest of the statements in that function will have access to that variable

Variable Scope

```
f <- function(x) {  
  y <- 1  
  g <- function(x) {  
    (x + y) / 2  
  }  
  g(x)  
}  
  
f(5)  
  
## [1] 3
```

`g()` is a subfunction that have access to `y` in `f`'s environment.

Variable Scope

```
f <- function(x) {  
  y <- 1  
  g(x)  
}
```

```
g <- function(x) {  
  (x + y) / 2  
}
```

```
f(5)
```

```
## Error in g(x): object 'y' not found
```

`g()` is a function that doesn't have access to `y`; `g()` can only see things in the global environment

One more thing ...

Let's look at another exmaple

```
mean(1:5)

## [1] 3

mean

## function (x, ...)
## UseMethod("mean")
## <bytecode: 0x7fe39d58e6d0>
## <environment: namespace:base>
```

One more thing ...

You can do things like this

```
# confusing but it works  
mean <- 1:5  
mean(mean)  
  
## [1] 3
```

Some issues

You can also do things like this

```
# not a good idea but you can do it  
mean <- function(x) 2*x + 5  
  
mean(1:5)  
  
## [1] 7 9 11 13 15
```

It seems we've lost the original mean() function

The :: Operator

:: operator to the rescue

```
# my mean  
mean(1:5)  
  
## [1]  7  9 11 13 15  
  
# base mean  
base::mean(1:5)  
  
## [1] 3
```

Here we use the name space base of the R package "base" to access the original mean()

Your Turn

Exercise

R has a function `summary()` that when applied on a numeric vector provides something like this:

```
summary(1:10)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.00	3.25	5.50	5.50	7.75	10.00

Create a `describe()` function that takes a numeric vector and returns: minimum, maximum, mean, and standard deviation

Exercise

First attempt

```
describe <- function(x) {  
  x_min <- min(x)  
  x_max <- max(x)  
  x_mean <- mean(x)  
  x_sd <- sd(x)  
  return(c(x_min, x_max, x_mean, x_sd))  
}  
  
describe(1:10)  
  
## [1] 1.00000 10.00000 7.00000 9.00000 11.00000 13.00000 15.  
## [9] 19.00000 21.00000 23.00000 25.00000 3.02765
```


Exercise

Second attempt (adding names)

```
describe <- function(x) {  
  x_min <- min(x)  
  x_max <- max(x)  
  x_mean <- mean(x)  
  x_sd <- sd(x)  
  values <- c(x_min, x_max, x_mean, x_sd)  
  names(values) <- c("min", "max", "mean", "sd")  
  return(values)  
}
```

```
describe(1:10)
```

```
##      min      max      mean      sd      <NA>      <NA>      <NA>  
## 1.00000 10.00000  7.00000  9.00000 11.00000 13.00000 15.00000  
##      <NA>      <NA>      <NA>      <NA>      <NA>  
## 19.00000 21.00000 23.00000 25.00000  3.02765
```

Exercise

Third attempt (using a list as output)

```
describe <- function(x) {  
  list(  
    min = min(x),  
    max = max(x),  
    mean = mean(x),  
    sd = sd(x)  
  )  
}
```

```
describe(1:10)  
  
## $min  
## [1] 1  
##  
## $max  
## [1] 10  
##  
## $mean  
## [1] 7 9 11 13 15 17 19 2  
##  
## $sd  
## [1] 3.02765
```

Exercise

Probability Density of the Normal Distribution:

$$f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Write a function that takes a value x (with parameters μ and σ) which computes the probability density distribution of the normal distribution

Exercise

Normal Distribution:

```
normal_dist <- function(x, mu = 0, sigma = 1) {  
  constant <- 1 / (sigma * sqrt(2*pi))  
  constant * exp(-((x - mu)^2) / (2 * sigma^2))  
}
```

```
normal_dist(2)
```

```
## [1] 0.05399097
```

Argument Matching

```
normal_dist <- function(x, mu = 0, sigma = 1) {  
  constant <- 1 / (sigma * sqrt(2*pi))  
  constant * exp(-((x - mu)^2) / (2 * sigma^2))  
}
```

```
normal_dist(2)  
normal_dist(2, sigma = 3, mu = 1)  
normal_dist(mu = 1, sigma = 3, 2)  
normal_dist(mu = 1, 2, sigma = 3)
```

Argument Matching

R is “smart” enough in doing pattern matching with arguments' names (not recommended though)

```
normal_dist(2)
```

```
## [1] 0.05399097
```

```
normal_dist(2, m = 0, s = 1)
```

```
## [1] 0.05399097
```

```
normal_dist(2, sig = 1, m = 0)
```

```
## [1] 0.05399097
```